

# Lecture 5: More on Kernels. SVM regression and classification

- More on kernels: How to tell if a function is a kernel?
- SVM classification
- SVM regression

## The “kernel trick”

- Recall: kernel functions are ways of expressing dot-products in some feature space:

$$K(\mathbf{x}, \mathbf{z}) = \phi(\mathbf{x}) \cdot \phi(\mathbf{z})$$

- In many cases,  $K$  can be computed in time that depends on the size of the inputs  $\mathbf{x}$  not the size of the feature space  $\phi(\mathbf{x})$
- If we work with a “dual” representation of the learning algorithm, we do not actually have to compute the feature mapping  $\phi$ . We just have to compute the similarity  $K$ .

## Polynomial kernels

- More generally,  $K(\mathbf{x}, \mathbf{z}) = (\mathbf{x} \cdot \mathbf{z})^d$  is a kernel, for any positive integer  $d$ :

$$K(\mathbf{x}, \mathbf{z}) = \left( \sum_{i=1}^n x_i z_i \right)^d$$

- If we expanded the sum above in the obvious way, we get  $n^d$  terms (i.e. feature expansion)
- Terms are monomials (products of  $x_i$ ) with total power equal to  $d$ .
- *Curse of dimensionality*: it is very expensive both to optimize and to predict in primal form
- However, *evaluating the dot-product of any two feature vectors can be done using  $K$  in  $O(n)$ !*

## Some other (fairly generic) kernel functions

- $K(\mathbf{x}, \mathbf{z}) = (1 + \mathbf{x} \cdot \mathbf{z})^d$  – feature expansion has all monomial terms of  $\leq d$  total power.
- Radial basis/Gaussian kernel:

$$K(\mathbf{x}, \mathbf{z}) = \exp(-\|\mathbf{x} - \mathbf{z}\|^2 / 2\sigma^2)$$

The kernel has an infinite-dimensional feature expansion, but dot-products can still be computed in  $O(n)$ !

- Sigmoidal kernel:

$$K(\mathbf{x}, \mathbf{z}) = \tanh(c_1 \mathbf{x} \cdot \mathbf{z} + c_2)$$

## Recall: Dual-view regression

- By re-writing the parameter vector as a linear combination of instances and solving, we get:

$$\mathbf{a} = (\mathbf{K} + \lambda \mathbf{I}_m)^{-1} \mathbf{y}$$

- The *feature mapping is not needed* either to learn or to make predictions!
- This approach is useful if the feature space is very large

## Making predictions in the dual view

- For a new input  $\mathbf{x}$ , the prediction is:

$$h(\mathbf{x}) = \mathbf{w}^\top \phi(\mathbf{x}) = \mathbf{a}^\top \Phi \phi(\mathbf{x}) = \mathbf{k}(\mathbf{x})^\top (\mathbf{K} + \lambda \mathbf{I}_m)^{-1} \mathbf{y}$$

where  $\mathbf{k}(\mathbf{x})$  is an  $m$ -dimensional vector, with the  $i$ th element equal to  $K(\mathbf{x}, \mathbf{x}_i)$

- That is, the  $i$ th element has the similarity of the input to the  $i$ th instance
- The features are not needed for this step either!
- This is a *non-parametric* representation - its size scales with the number of instances.

# Kernels

- A lot of current research has to do with defining new kernels functions, suitable to particular tasks / kinds of input objects
- Many kernels are available:
  - Information diffusion kernels (Lafferty and Lebanon, 2002)
  - Diffusion kernels on graphs (Kondor and Jebara 2003)
  - String kernels for text classification (Lodhi et al, 2002)
  - String kernels for protein classification (e.g., Leslie et al, 2002)
  - ... and others!

## Example: String kernels

- Very important for DNA matching, text classification, ...
- Example: in DNA matching, we use a sliding window of length  $k$  over the two strings that we want to compare
- The window is of a given size, and inside we can do various things:
  - Count exact matches
  - Weigh mismatches based on how bad they are
  - Count certain markers, e.g. AGT
- The kernel is the sum of these similarities over the two sequences
- How do we prove this is a kernel?



## Establishing “kernelhood”

- Suppose someone hands you a function  $K$ . How do you know that it is a kernel?
- More precisely, given a function  $K : \mathbb{R}^n \times \mathbb{R}^n \rightarrow \mathbb{R}$ , under what conditions can  $K(\mathbf{x}, \mathbf{z})$  be written as a dot product  $\phi(\mathbf{x}) \cdot \phi(\mathbf{z})$  for some feature mapping  $\phi$ ?
- We want a general recipe, which does not require explicitly defining  $\phi$  every time

## Kernel matrix

- Suppose we have an arbitrary set of input vectors  $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_m$
- The *kernel matrix (or Gram matrix)*  $\mathbf{K}$  corresponding to kernel function  $K$  is an  $m \times m$  matrix such that  $\mathbf{K}_{ij} = K(\mathbf{x}_i, \mathbf{x}_j)$  (notation is overloaded on purpose).
- What properties does the kernel matrix  $\mathbf{K}$  have?
- Claims:
  1.  $\mathbf{K}$  is symmetric
  2.  $\mathbf{K}$  is positive semidefinite
- Note that these claims are consistent with the intuition that  $K$  is a “similarity” measure (and will be true regardless of the data)

## Proving the first claim

If  $K$  is a valid kernel, then the kernel matrix is symmetric

$$\mathbf{K}_{ij} = \phi(\mathbf{x}_i) \cdot \phi(\mathbf{x}_j) = \phi(\mathbf{x}_j) \cdot \phi(\mathbf{x}_i) = \mathbf{K}_{ji}$$

## Proving the second claim

If  $K$  is a valid kernel, then the kernel matrix is positive semidefinite

Proof: Consider an arbitrary vector  $\mathbf{z}$

$$\begin{aligned}\mathbf{z}^\top \mathbf{K} \mathbf{z} &= \sum_i \sum_j z_i K_{ij} z_j = \sum_i \sum_j z_i (\phi(\mathbf{x}_i) \cdot \phi(\mathbf{x}_j)) z_j \\ &= \sum_i \sum_j z_i \left( \sum_k \phi_k(\mathbf{x}_i) \phi_k(\mathbf{x}_j) \right) z_j \\ &= \sum_k \sum_i \sum_j z_i \phi_k(\mathbf{x}_i) \phi_k(\mathbf{x}_j) z_j \\ &= \sum_k \left( \sum_i z_i \phi_k(\mathbf{x}_i) \right)^2 \geq 0\end{aligned}$$

## Mercer's theorem

- We have shown that if  $K$  is a kernel function, then for any data set, the corresponding kernel matrix  $\mathbf{K}$  defined by  $\mathbf{K}_{ij} = K(\mathbf{x}_i, \mathbf{x}_j)$  is symmetric and positive semidefinite
- Mercer's theorem states that the reverse is also true:  
Given a function  $K : \mathbb{R}^n \times \mathbb{R}^n \rightarrow \mathbb{R}$ ,  *$K$  is a kernel if and only if, for any data set, the corresponding kernel matrix is symmetric and positive semidefinite*
- The reverse direction of the proof is much harder (see e.g. Vapnik's book for details)
- This result gives us a way to check if a given function is a kernel, by checking these two properties of its kernel matrix.
- Kernels can also be obtained by combining other kernels, or by learning from data
- Kernel learning may suffer from overfitting (kernel matrix close to diagonal)

## More on RKHS

- Mercer's theorem tells us that a function  $K : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$  is a kernel (i.e. it corresponds to the inner product in some feature space) if and only if it is positive semi-definite.
- The feature space is the **reproducing kernel Hilbert space** (RKHS)

$$\mathcal{H} = \left\{ \sum_j \alpha_j K(z_j, \cdot) : z_j \in \mathcal{X}, \alpha_j \in \mathbb{R} \right\}$$

with inner product  $\langle K(z, \cdot), K(z', \cdot) \rangle_{\mathcal{H}} = K(z, z')$ .

- The term reproducing comes from the **reproducing property** of the kernel function:

$$\forall f \in \mathcal{H}, x \in \mathcal{X} : f(x) = \langle f(\cdot), K(x, \cdot) \rangle_{\mathcal{H}}$$

- Recall that the solution of the regularized least square in the feature space associated to a kernel function  $K$  has the form  $h(\mathbf{x}) = \sum_{i=1}^m \alpha_i K(\mathbf{x}_i, \mathbf{x})$ . This is a particular case of the **representer theorem**...

## Representer Theorem

**Theorem 1.** *Let  $K : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$  be a positive definite kernel and let  $\mathcal{H}$  be the corresponding RKHS.*

*Then for any training sample  $\mathcal{S} = \{(x_i, y_i)\}_{i=1}^m \subset \mathcal{X} \times \mathbb{R}$ , any loss function  $\ell : (\mathcal{X} \times \mathbb{R} \times \mathbb{R})^m \rightarrow \mathbb{R}$  and any real-valued non-decreasing function  $g$ , the solution of the optimization problem*

$$\arg \min_{f \in \mathcal{H}} \ell((x_1, y_1, f(x_1)), \dots, (x_m, y_m, f(x_m))) + g(\|f\|_{\mathcal{H}})$$

*admits a representation of the form*

$$f^*(\cdot) = \sum_{i=1}^m \alpha_i K(\mathbf{x}_i, \cdot).$$

[Schölkopf, Herbrich and Smola. *A generalized representer Theorem*. COLT 2001.]

## Regularization in the dual view

- We want to penalize the function we are trying to estimate (to keep it simple)
- Assume this is part of a reproducing kernel Hilbert space
- We want to minimize:


$$J(h) = \frac{1}{2} \sum_{i=1}^n (y_i - h(\mathbf{x}_i))^2 + \frac{\lambda}{2} \|h\|_{\mathcal{H}}^2$$

- If we put a Gaussian prior on  $h$ , and solve, we obtain *Gaussian process regression*



## Binary classification revisited


- Consider a linearly separable binary classification data set  $\{\mathbf{x}_i, y_i\}_{i=1}^m$ .
- There is an infinite number of hyperplanes that separate the classes:



- Which plane is best?
- Relatedly, for a given plane, for which points should we be most confident in the classification?


## The margin, and linear SVMs

- For a given separating hyperplane, the *margin* is two times the (Euclidean) distance from the hyperplane to the nearest training example.



- It is the width of the “strip” around the decision boundary containing no training examples.
- A linear SVM is a perceptron for which we choose  $\mathbf{w}, w_0$  so that margin is maximized

# Hyperplanes and scale invariance




- **Hyperplane:**  $\mathbf{w}^\top \mathbf{x} + b = 0$ .
- $\Rightarrow$  Scale invariance:  $(\alpha \mathbf{w})^\top \mathbf{x} + \alpha b = 0$  define the same hyperplane for any scalar  $\alpha \in \mathbb{R}$ .

Figures taken from David Sontag's lecture slides:

<http://people.csail.mit.edu/dsontag/courses/ml13/slides/lecture3.pdf>

# Hyperplanes and scale invariance



- **Hyperplane:**  $\mathbf{w}^\top \mathbf{x} + b = 0$ .
  - ⇒ Scale invariance:  $\alpha \mathbf{w}^\top \mathbf{x} + \alpha b = 0 \dots$
  - Workaround: we set the scale by enforcing
$$\mathbf{w}^\top \mathbf{x} + b \geq 1 \quad \text{for all positive examples}$$
$$\mathbf{w}^\top \mathbf{x} + b \leq -1 \quad \text{for all negative examples}$$
- equivalently, we want to satisfy the **linear constraints**

$$y_i(\mathbf{w}^\top \mathbf{x} + b) \geq 1$$

for all  $i$  (here  $\mathcal{Y} = \{-1, 1\}$ ).

## Margin as a function of the normal vector

- Let  $\mathbf{x}_1$  be a point on the boundary and  $\mathbf{x}_2$  its projection on the hyperplane:

$$\mathbf{w}^\top \mathbf{x}_1 + b = 1$$


$$\mathbf{w}^\top \mathbf{x}_2 + b = 0$$

We get  $\mathbf{w}^\top (\mathbf{x}_1 - \mathbf{x}_2) = 1$ .

- But we also know that

$$\mathbf{x}_1 - \mathbf{x}_2 = \gamma \frac{\mathbf{w}}{\|\mathbf{w}\|}$$

- Hence  $\gamma = \frac{1}{\|\mathbf{w}\|}$ : maximizing the margin is equivalent to **minimizing the norm of  $\mathbf{w}$** !



## SVM final formulation

- Let's minimize  $\|\mathbf{w}\|^2$  instead of  $\|\mathbf{w}\|$ .  
(it will be easier and taking the square is a monotone transformation, as  $\|\mathbf{w}\|$  is positive, so this doesn't change the optimal solution.)

- This gets us to:



$$\begin{array}{ll}\min & \|\mathbf{w}\|^2 \\ \text{w.r.t.} & \mathbf{w}, w_0 \\ \text{s.t.} & y_i(\mathbf{w} \cdot \mathbf{x}_i + w_0) \geq 1\end{array}$$

- This we can solve! How?
  - It is a *quadratic programming* (QP) problem—a standard type of optimization problem<sup>1</sup> for which many efficient packages are available.
  - Better yet, it's a convex (positive semidefinite) QP

---

<sup>1</sup>optimization of a quadratic function of several variables subject to linear constraints on these variables.

## Example



We have a solution, but no support vectors yet...

## Lagrange multipliers for inequality constraints (revisited)

- Suppose we have the following optimization problem, called *primal*:

$$\begin{aligned} & \min_{\mathbf{w}} f(\mathbf{w}) \\ & \text{such that } g_i(\mathbf{w}) \leq 0, \quad i = 1 \dots k \end{aligned}$$

- We define the *Lagrangian*:

$$L(\mathbf{w}, \boldsymbol{\alpha}) = f(\mathbf{w}) + \sum_{i=1}^k \alpha_i g_i(\mathbf{w}), \quad (1)$$

where  $\alpha_i$ ,  $i = 1 \dots k$  are the Lagrange multipliers.



## A different optimization problem

$$L(\mathbf{w}, \alpha) = f(\mathbf{w}) + \sum_{i=1}^k \alpha_i g_i(\mathbf{w})$$

- Consider  $\mathcal{P}(\mathbf{w}) = \max_{\alpha: \alpha_i \geq 0} L(\mathbf{w}, \alpha)$
- Observe that the following is true:

$$\mathcal{P}(\mathbf{w}) = \begin{cases} f(\mathbf{w}) & \text{if all constraints are satisfied} \\ +\infty & \text{otherwise} \end{cases}$$

- Hence, instead of computing  $\min_{\mathbf{w}} f(\mathbf{w})$  subject to the original constraints, we can compute:

$$p^* = \min_{\mathbf{w}} \mathcal{P}(\mathbf{w}) = \min_{\mathbf{w}} \max_{\alpha: \alpha_i \geq 0} L(\mathbf{w}, \alpha)$$

## Dual optimization problem

- We defined the **primal problem**:
- The **dual problem** is obtained by swapping min and max:

$$p^* = \min_{\mathbf{w}} \max_{\alpha: \alpha_i \geq 0} L(\mathbf{w}, \alpha)$$

$$d^* = \max_{\alpha: \alpha_i \geq 0} \min_{\mathbf{w}} L(\mathbf{w}, \alpha)$$

- We can show that  $d^* \leq p^*$ :
  - Let  $p^* = L(w^p, \alpha^p)$ , let  $d^* = L(w^d, \alpha^d)$ .
  - Then  $d^* = L(w^d, \alpha^d) \leq L(w^p, \alpha^d) \leq L(w^p, \alpha^p) = p^*$ .
- The non-negative number  $p^* - d^*$  is called the **duality gap**.

## Dual optimization problem (cont'd)

- **Slater's condition**: if  $f$  (objective) and the  $g_i$ 's (constraints) are convex, and the  $g_i$ 's can all be satisfied simultaneously for some  $\mathbf{w}$ , then we have equality:  $d^* = p^* = L(\mathbf{w}^*, \alpha^*)$ .
- Which means that **the two problems are equivalent!**
- Moreover  $\mathbf{w}^*, \alpha^*$  solve the primal and dual if and only if they satisfy the following conditions (called Karush-Kuhn-Tucker):

$$\frac{\partial}{\partial w_i} L(\mathbf{w}^*, \alpha^*) = 0, \quad i = 1 \dots n \quad (2)$$

$$\alpha_i^* g_i(\mathbf{w}^*) = 0, \quad i = 1 \dots k \quad (3)$$

$$g_i(\mathbf{w}^*) \leq 0, \quad i = 1 \dots k \quad (4)$$

$$\alpha_i^* \geq 0, \quad i = 1 \dots k \quad (5)$$

## Back to maximum margin perceptron

- We wanted to solve (rewritten slightly):

$$\begin{array}{ll}\min & \frac{1}{2} \|\mathbf{w}\|^2 \\ \text{w.r.t.} & \mathbf{w}, w_0 \\ \text{s.t.} & 1 - y_i(\mathbf{w} \cdot \mathbf{x}_i + w_0) \leq 0\end{array}$$

- The Lagrangian is:

$$L(\mathbf{w}, w_0, \alpha) = \frac{1}{2} \|\mathbf{w}\|^2 + \sum_i \alpha_i (1 - y_i(\mathbf{w} \cdot \mathbf{x}_i + w_0))$$

- The primal problem is:  $\min_{\mathbf{w}, w_0} \max_{\alpha: \alpha_i \geq 0} L(\mathbf{w}, w_0, \alpha)$
- We will solve the dual problem:  $\max_{\alpha: \alpha_i \geq 0} \min_{\mathbf{w}, w_0} L(\mathbf{w}, w_0, \alpha)$
- In this case, the optimal solutions coincide, because we have a quadratic objective and linear constraints (both of which are convex).

## Solving the dual

- Lagrangian:  $L(\mathbf{w}, w_0, \alpha) = \frac{1}{2}\|\mathbf{w}\|^2 + \sum_i \alpha_i(1 - y_i(\mathbf{w} \cdot \mathbf{x}_i + w_0))$
- From KKT (2), the derivatives of  $L(\mathbf{w}, w_0, \alpha)$  wrt  $\mathbf{w}, w_0$  should be 0:

$$\frac{\partial L}{\partial w_0} = -\sum_i \alpha_i y_i \quad \Rightarrow \quad \sum_i \alpha_i y_i = 0$$

$$\nabla_{\mathbf{w}} L = \mathbf{w} - \sum_i \alpha_i y_i \mathbf{x}_i \quad \Rightarrow \quad \mathbf{w} = \sum_i \alpha_i y_i \mathbf{x}_i$$

- Observe that  $\alpha_i > 0$  implies that the corresponding constraint is tight, i.e.  $y_i(\mathbf{w} \cdot \mathbf{x}_i + w_0) = 1$ , from which we obtain

$$w_0 = y_i - \mathbf{w} \cdot \mathbf{x}_i \quad \text{for all } i \text{ such that } \alpha_i > 0.$$

$\Rightarrow$  Just like for the perceptron with zero initial weights, the optimal solution for  $\mathbf{w}$  is a linear combination of the  $\mathbf{x}_i$ .

- The output is

$$h_{\mathbf{w}, w_0}(\mathbf{x}) = \text{sign} \left( \sum_{i=1}^m \alpha_i y_i (\mathbf{x}_i \cdot \mathbf{x}) + w_0 \right)$$

⇒ Output depends on weighted dot product of input vector with training examples

## Solving the dual (cont'd)

- By plugging these back into the expression for  $L$ , we get:

$$\max_{\alpha} \sum_i \alpha_i - \frac{1}{2} \sum_{i,j} y_i y_j \alpha_i \alpha_j (\mathbf{x}_i \cdot \mathbf{x}_j)$$

with constraints:  $\alpha_i \geq 0$  and  $\sum_i \alpha_i y_i = 0$

## The support vectors


- Suppose we find optimal  $\alpha$ s (e.g., using a standard QP package)
- The  $\alpha_i$  will be  $> 0$  only for the points for which  $1 - y_i(\mathbf{w} \cdot \mathbf{x}_i + w_0) = 0$
- These are the points lying on the edge of the margin, and they are called *support vectors*, because they define the decision boundary
- The output of the classifier for query point  $\mathbf{x}$  is computed as:

$$\text{sgn} \left( \sum_{i=1}^m \alpha_i y_i (\mathbf{x}_i \cdot \mathbf{x}) + w_0 \right)$$

Hence, the output is determined by computing the *dot product of the point with the support vectors*!




## Example



Support vectors are in bold

## Non-linearly separable data



- A linear boundary might be too simple to capture the class structure.
- One way of getting a nonlinear decision boundary in the input space is to find a linear decision boundary in an expanded space (e.g., for polynomial regression.)
- Thus,  $\mathbf{x}_i$  is replaced by  $\phi(\mathbf{x}_i)$ , where  $\phi$  is called a *feature mapping*

## Margin optimization in feature space

- Replacing  $\mathbf{x}_i$  with  $\phi(\mathbf{x}_i)$ , the optimization problem to find  $\mathbf{w}$  and  $w_0$  becomes:

$$\begin{array}{ll}\min & \|\mathbf{w}\|^2 \\ \text{w.r.t.} & \mathbf{w}, w_0 \\ \text{s.t.} & y_i(\mathbf{w} \cdot \phi(\mathbf{x}_i) + w_0) \geq 1\end{array}$$

- Dual form:

$$\begin{array}{ll}\max & \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i,j=1}^m y_i y_j \alpha_i \alpha_j \phi(\mathbf{x}_i) \cdot \phi(\mathbf{x}_j) \\ \text{w.r.t.} & \alpha_i \\ \text{s.t.} & 0 \leq \alpha_i \\ & \sum_{i=1}^m \alpha_i y_i = 0\end{array}$$

$\Rightarrow$  We just need to compute  $\phi(\mathbf{x}_i) \cdot \phi(\mathbf{x}_j)$  for each  $i, j$  to solve the dual problem, never need to compute the feature map explicitly...

## Feature space solution

- The optimal weights, in the expanded feature space, are given by

$$\mathbf{w} = \sum_{i=1}^m \alpha_i y_i \phi(\mathbf{x}_i) \quad \text{and} \quad w_0 = y_j - \sum_{i=1}^m \alpha_i y_i \phi(\mathbf{x}_i) \cdot \phi(\mathbf{x}_j)$$

where  $j$  is such that  $\alpha_j > 0$ .

- Classification of an input  $\mathbf{x}$  is given by:

$$h_{\mathbf{w}, w_0}(\mathbf{x}) = \text{sign} \left( \sum_{i=1}^m \alpha_i y_i \phi(\mathbf{x}_i) \cdot \phi(\mathbf{x}) + w_0 \right)$$

⇒ To solve the SVM optimization problem in dual form and to make a prediction, *we only need to compute dot-products of feature vectors*.

⇒ SVMs can straightforwardly be kernelized!

## The “kernel trick”

- If we work with the dual, we do not actually have to ever compute the feature mapping  $\phi$ . We just have to compute the similarity  $K$ .
- That is, we can solve the dual for the  $\alpha_i$ :

$$\begin{array}{ll}\max & \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i,j=1}^m y_i y_j \alpha_i \alpha_j K(\mathbf{x}_i, \mathbf{x}_j) \\ \text{w.r.t.} & \alpha_i \\ \text{s.t.} & 0 \leq \alpha_i \\ & \sum_{i=1}^m \alpha_i y_i = 0\end{array}$$

- The class of a new input  $\mathbf{x}$  is computed as:

$$h_{\mathbf{w}, w_0}(\mathbf{x}) = \text{sign} \left( \left( \sum_{i=1}^m \alpha_i y_i \phi(\mathbf{x}_i) \right) \cdot \phi(\mathbf{x}) + w_0 \right) = \text{sign} \left( \sum_{i=1}^m \alpha_i y_i K(\mathbf{x}_i, \mathbf{x}) + w_0 \right)$$

## Regularization with SVMs

- Kernels are a powerful tool for allowing non-linear, complex functions
- But now the number of parameters can be as high as the number of instances!
- With a very specific, non-linear kernel, each data point may be in its own partition
- With linear and logistic regression, we used regularization to avoid overfitting
- We need a method for allowing regularization with SVMs as well.

## Soft margin SVM

- Recall that in the linearly separable case, we compute the solution to the following optimization problem:

$$\begin{array}{ll}\min & \frac{1}{2} \|\mathbf{w}\|^2 \\ \text{w.r.t.} & \mathbf{w}, w_0 \\ \text{s.t.} & y_i(\mathbf{w} \cdot \mathbf{x}_i + w_0) \geq 1\end{array}$$

- If we want to allow misclassifications, we can relax the constraints to:

$$y_i(\mathbf{w} \cdot \mathbf{x}_i + w_0) \geq 1 - \xi_i$$

- If  $\xi_i \in (0, 1)$ , the data point is within the margin
- If  $\xi_i \geq 1$ , then the data point is misclassified
- We define the **soft error** as  $\sum_i \xi_i$
- We will have to change the criterion to reflect the soft errors

## New problem formulation with soft errors

Instead of:

$$\begin{array}{ll}\min & \frac{1}{2} \|\mathbf{w}\|^2 \\ \text{w.r.t.} & \mathbf{w}, w_0 \\ \text{s.t.} & y_i(\mathbf{w} \cdot \mathbf{x}_i + w_0) \geq 1\end{array}$$

we want to solve:

$$\begin{array}{ll}\min & \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_i \xi_i \\ \text{w.r.t.} & \mathbf{w}, w_0, \xi_i \\ \text{s.t.} & y_i(\mathbf{w} \cdot \mathbf{x}_i + w_0) \geq 1 - \xi_i, \xi_i \geq 0\end{array}$$

- Note that soft errors include points that are misclassified, as well as points within the margin
- There is a linear penalty for both categories
- The choice of the *constant  $C$  controls overfitting*



## A built-in overfitting knob

$$\begin{array}{ll}\min & \frac{1}{2}\|\mathbf{w}\|^2 + C \sum_i \xi_i \\ \text{w.r.t.} & \mathbf{w}, w_0, \xi_i \\ \text{s.t.} & y_i(\mathbf{w} \cdot \mathbf{x}_i + w_0) \geq 1 - \xi_i \\ & \xi_i \geq 0\end{array}$$

- If  $C$  is 0, there is no penalty for soft errors, so the focus is on maximizing the margin, even if this means more mistakes
- If  $C$  is very large, the emphasis on the soft errors will cause decreasing the margin, if this helps to classify more examples correctly.
- Internal cross-validation is a good way to choose  $C$  appropriately

## Lagrangian for the new problem

- Like before, we can write a Lagrangian for the problem and then use the dual formulation to find the optimal parameters:

$$\begin{aligned} L(\mathbf{w}, w_0, \alpha, \xi, \mu) &= \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_i \xi_i \\ &+ \sum_i \alpha_i (1 - \xi_i - y_i(\mathbf{w}_i \cdot \mathbf{x}_i + w_0)) + \sum_i \mu_i \xi_i \end{aligned}$$

- All the previously described machinery can be used to solve this problem
- Note that in addition to  $\alpha_i$  we have coefficients  $\mu_i$ , which ensure that the errors are positive, but do not participate in the decision boundary

## Soft margin optimization with kernels

- Replacing  $\mathbf{x}_i$  with  $\phi(\mathbf{x}_i)$ , the optimization problem to find  $\mathbf{w}$  and  $w_0$  becomes:

$$\begin{array}{ll}\min & \|\mathbf{w}\|^2 + C \sum_i \xi_i \\ \text{w.r.t.} & \mathbf{w}, w_0, \xi \\ \text{s.t.} & y_i(\mathbf{w} \cdot \phi(\mathbf{x}_i) + w_0) \geq (1 - \xi_i) \\ & \xi_i \geq 0\end{array}$$

- Dual form and solution have similar forms to what we described last time, but in terms of kernels

## Getting SVMs to work in practice

- Two important choices:
  - Kernel (and kernel parameters)
  - Regularization parameter  $C$
- The parameters may interact!  
E.g. for Gaussian kernel, the larger the width of the kernel, the more biased the classifier, so low  $C$  is better
- Together, these control overfitting: always do an internal parameter search, using a validation set!
- Overfitting symptoms:
  - Low margin
  - Large fraction of instances are support vectors


## Solving the quadratic optimization problem

- Many approaches exist
- Because we have constraints, gradient descent does not apply directly (the optimum might be outside of the feasible region)
- Platt's algorithm is the fastest current approach, based on *coordinate ascent*

## Coordinate ascent

- Suppose you want to find the maximum of some function  $F(\alpha_1, \dots, \alpha_n)$  (the  $\alpha_i$  could be scalars, vectors, matrices...)
- Coordinate ascent optimizes the function by repeatedly picking an  $\alpha_i$  and optimizing it, while all other parameters are fixed:
  1. Initialize the  $\alpha_i$ 's
  2. Repeat until convergence (or stopping criterion):
    - Pick an  $i$  (round-robin, random, ...)
    - Solve  $\alpha_i \leftarrow \arg \min_{\alpha_i} F(\alpha_1, \dots, \alpha_n)$
- There are different ways of looping through the parameters:
  - Round-robin
  - Repeatedly pick a parameter at random
  - Choose next the variable expected to make the largest improvement
  - ...

## Example



## Our optimization problem (dual form)

$$\max_{\alpha} \sum_i \alpha_i - \frac{1}{2} \sum_{i,j} y_i y_j \alpha_i \alpha_j (\phi(\mathbf{x}_i) \cdot \phi(\mathbf{x}_j))$$

with constraints:  $0 \leq \alpha_i \leq C$  and  $\sum_i \alpha_i y_i = 0$

- Suppose we want to optimize for  $\alpha_1$  while  $\alpha_2, \dots, \alpha_n$  are fixed
- We cannot do it because  $\alpha_1$  will be completely determined by the last constraint:  $\alpha_1 = -y_1 \sum_{i=2}^m \alpha_i y_i$
- Instead, we have to optimize *pairs of parameters*  $\alpha_i, \alpha_j$  together




## SMO

- Suppose that we want to optimize  $\alpha_1$  and  $\alpha_2$  together, while all other parameters are fixed.
- We know that  $y_1\alpha_1 + y_2\alpha_2 = -\sum_{i=1}^m y_i\alpha_i = \xi$ , where  $\xi$  is a constant
- So  $\alpha_1 = y_1(\xi - y_2\alpha_2)$  (because  $y_1$  is either  $+1$  or  $-1$  so  $y_1^2 = 1$ )
- This defines a line, and any pair  $\alpha_1, \alpha_2$  which is a solution has to be on the line

## SMO (II)

- We also know that  $0 \leq \alpha_1 \leq C$  and  $0 \leq \alpha_2 \leq C$ , so the solution has to be on the line segment inside the rectangle below



## SMO(III)

- By plugging  $\alpha_1$  back in the optimization criterion, we obtain a quadratic function of  $\alpha_2$ , whose optimum we can find exactly
- If the optimum is inside the rectangle, we take it.
- If not, we pick the closest intersection point of the line and the rectangle
- This procedure is very fast because all these are simple computations.

# Complexity

- Quadratic programming is expensive in the number of training examples
- Platt's SMO algorithm is quite fast though, and other fancy optimization approaches are available
- Best packages can handle 50,000+ instances, but not more than 100,000
- On the other hand, number of attributes can be very high (strength compared to neural nets)
- Evaluating a SVM is *slow if there are a lot of support vectors*.
- Dictionary methods attempt to select a subset of the data on which to train.

# Support Vector Regression

- In regression problems, so far we have been trying to minimize mean-squared error:

$$\sum_i (y_i - (\mathbf{w} \cdot \mathbf{x}_i + w_0))^2$$

- In SVM regression, we will be interested instead in minimizing absolute error:

$$\sum_i |y_i - (\mathbf{w} \cdot \mathbf{x}_i + w_0)|$$


- This is more *robust to outliers* than the squared loss
- But we cannot require that all points be approximated correctly (overfitting!)

## Loss function for support vector regression

In order to allow for misclassifications in SVM regression (and have robustness to noise), we use the  *$\varepsilon$ -insensitive loss*:

$$J_{\varepsilon} = \sum_{i=1}^m J_{\varepsilon}(\mathbf{x}_i), \text{ where}$$

$$J_{\varepsilon}(\mathbf{x}_i) = \begin{cases} 0 & \text{if } |y_i - (\mathbf{w} \cdot \mathbf{x}_i + w_0)| \leq \varepsilon \\ |y_i - (\mathbf{w} \cdot \mathbf{x}_i + w_0)| - \varepsilon & \text{otherwise} \end{cases}$$



## Solving SVM regression

- We introduce slack variables,  $\xi_i^+$ ,  $\xi_i^-$  to account for errors outside the tolerance area
- We need two kinds of variables to account for both positive and negative errors




## The optimization problem

$$\begin{array}{ll}\min & \frac{1}{2}\|\mathbf{w}\|^2 + C \sum_i (\xi_i^+ + \xi_i^-) \\ \text{w.r.t.} & \mathbf{w}, w_0, \xi_i^+, \xi_i^- \\ \text{s.t.} & y_i - (\mathbf{w} \cdot \mathbf{x}_i + w_0) \leq \varepsilon + \xi_i^+ \\ & y_i - (\mathbf{w} \cdot \mathbf{x}_i + w_0) \geq -\varepsilon - \xi_i^- \\ & \xi_i^+, \xi_i^- \geq 0\end{array}$$

- Like before, we can write the Lagrangian and solve the dual form of the problem
- Kernels can be used as before to get non-linear functions



## Effect of $\varepsilon$



3

- As  $\varepsilon$  increases, the function is allowed to move away from the data points, the number of support vectors decreases and the fit gets worse

<sup>3</sup>Zisserman course notes